

# Economic Denial of Sustainability Attacks on Kubernetes





JONATHAN CHAMBERLAIN

PRESENTATION FOR SEMINAR ON PRACTICAL SECURITY 2023-08-10

# Overview

- ▶ The cloud really is “someone else’s computer”
  - ▶ Business model: server providers concern themselves with maintaining hardware, application owners need only focus on the software.
  - ▶ Allows for rapid capacity expansion: lease additional compute nodes as needed.
  - ▶ How to determine when to scale?

# Overview




# Overview

- ▶ Autoscaling is useful for DDoS defense.
  - ▶ Having autoscaling enabled is considered a security best practice by Amazon and others for this very reason [1].
  - ▶ Compute nodes are not free.
  - ▶ Thus, the defense is itself open to exploitation.
- ▶ This can result in an Economic Denial of Sustainability (EDoS) scenario.

# Overview

- ▶ EDoS in this context is accomplished by flooding the system with excess traffic to forcibly trigger the autoscaling mechanisms.






# Yo-Yo Attacks in Kubernetes

- ▶ A particular class of EDoS attack, the Yo-Yo, was demonstrated as a proof of concept by Ben David and Bremner-Barr. [2]
- ▶ The key finding is that with Yo-Yo, nearly the same economic cost is imposed on the target without the need to send continuous traffic.


# Yo-Yo Attacks in Kubernetes

- Consider an example DDoS attack on Kubernetes infrastructure [2]



# Yo-Yo Attacks in Kubernetes

- Contrast with the comparable Yo-Yo attack [2]





# Yo-Yo Attacks in Kubernetes

- ▶ While these experiments were done in a controlled environment, the underlying attack structure is increasingly common.
  - ▶ In 2021, wave attacks accounted for approximately half of all DoS style attacks [3]
  - ▶ However, a sample of corporate attendees at a Red Hat Research Days talk indicated that they had no plan for dealing with such attacks [4].


# Yo-Yo Attacks in Kubernetes

- ▶ Further, there is a question as to whether the Yo-Yo is really the “best” that can be done:
  - ▶ The attack as structured in [2] is not randomized in any way
  - ▶ In [2], scaling only occurs as a result of adversarial traffic
  - ▶ Attack/rest periods determined solely by scaling time assuming traffic goes from some constant level of normal traffic, to maximum adversarial traffic, back to normal traffic levels again


# Improving on the Yo-Yo Attack

- ▶ By default, the Cluster Autoscaler software checks every 10 seconds to determine whether the number of nodes need to be reallocated [5].
- ▶ Our hypothesis: an attacker can exploit timings/system states to inject just enough traffic to keep the system at a higher number of nodes for cheaper than the wave attack.

# Improving on the Yo-Yo Attack



# Improving on the Yo-Yo Attack





# Improving on the Yo-Yo Attack

- ▶ In addition, while the Cluster Autoscaler is itself open sourced on GitHub, the current deployment is through managed cloud providers [5].
  - ▶ Thus, certain settings become a “black box” which cannot be updated by the application owner.
- ▶ Our current efforts are focused on two areas:
  - ▶ Developing a Game Theoretic model
  - ▶ Running experiments on OpenShift utilizing the NERC

# Game Theoretic Approach

- ▶ Essentially, want a quantifiable and provably optimal policy which can be taken by the attacker, given assumptions on the system settings.
  - ▶ Allows for evaluation of how system changes impact attacker strategy
  - ▶ Can design defense around expected attacker behavior.
- ▶ Accomplish this through Markov Decision Processes (MDPs)

# Game Theoretic Approach

- ▶ MDPs are described by:
  - ▶ States, in this case  $(m,n)$  where  $m$  is the autoscaling metric of interest, and  $n$  is the number of active nodes
  - ▶ Actions, in this case to add a node, remove a node, or keep the number of nodes
  - ▶ Costs/rewards, e.g. cost of running extra nodes, cost of activating/deactivating nodes, refunds to customers for SLA violations, perceived costs of over/under provisioning

# Game Theoretic Approach

- ▶ At each step, the action which maximizes reward/minimizes cost is taken.
- ▶ Over time, the actions taken define a policy, which can be expressed as a Markov chain.
- ▶ In our case, this will yield the thresholds for activating/deactivating the next node.
  - ▶ If the attacker can glean the state of the system at any given time, this information allows for better attack planning.
- ▶ Solving MDPs is complex – currently adapting python code to accomplish this for our case(s).

# Experiments on OpenShift

- ▶ In addition and in parallel to the theoretic approach, we intend to run experiments on a NERC OpenShift instance through creating a cluster of our own.
- ▶ Ideally, this enables full control over the relevant settings, rather than some settings be confined to provider fiat, as in [2]



# Experiments on OpenShift

- In particular, the Node scale-up/scale-down can be controlled through Cluster Autoscaler [5]

Parameter	Definition	Configuration given by	Value
$r$	Average requests rate per second of legitimate clients	System usage	
$N_p$	Initial number of Pods	System administrator	4
$N_n$	Initial number of Nodes		4
$R$	Number of Pods per Node		3
$I_{up}^p \setminus I_{down}^p$	Threshold interval for scale-up and scale-down for a Pod		1min\5min
$I_{up}^n \setminus I_{down}^n$	Threshold interval for scale-up and scale-down for a Node		10sec\10min
$W_{up}^p \setminus W_{down}^p$	Warming time of scale-up and scale-down for a Pod	Kubernetes infrastructure	30sec\5sec
$W_{up}^n \setminus W_{down}^n$	Warming time of scale-up and scale-down for a Node		2min\2min
$k$	The power of the attack	Attacker	10\20
$n$	Number of attack cycles		
$T$	Cycle duration		
$t_{on} \setminus t_{off}$	Time of <i>on-attack</i> phase and <i>off-attack</i> phase. $T = t_{on} + t_{off}$		

# Experiments on OpenShift

- ▶ The main roadblock at present is that OpenShift magnum, which will allow for integration with Cluster Autoscaler, is not currently production code
  - ▶ Scheduled for October release as part of OpenShift 2023.2 update
- ▶ As a result, we are currently working to try and simulate cluster behavior via other means, as well as running experiments with minikube to determine behavior when the number of nodes is fixed.

# Conclusions

- ▶ Wave (a/k/a burst) attacks are an emerging threat to cloud infrastructure.
- ▶ Tools utilized to prevent such attacks from creating a DDoS scenario can be weaponized to create an EDoS scenario instead.
- ▶ Proof of concept in the form of periodic Yo-Yo attacks exists.
- ▶ Arguably not the most sophisticated attack possible, current research efforts are working towards proving this.
  - ▶ Goal is ultimately to develop better detection and defense capabilities against such attacks.

# References

- ▶ [1] Amazon Web Services, "AWS Best Practices for DDoS Resiliency," 2021.  
[https://d0.awsstatic.com/whitepapers/Security/DDoS White Paper.pdf](https://d0.awsstatic.com/whitepapers/Security/DDoS%20White%20Paper.pdf)
- ▶ [2] R. Ben David, and A. Bremler Barr. "Kubernetes autoscaling: Yoyo attack vulnerability and mitigation." *arXiv preprint arXiv:2105.00542* (2021).
- ▶ [3] Imperva. *2021 Cyberdefense Threat Report*. 2021.  
<https://www.imperva.com/resources/resource-library/reports/2021-cyberthreat-defense-report/>
- ▶ [4] A. Bremler-Barr and M. Czeizler. *Research Days 2023: Cloud Auto-scaling Mechanism Under DDoS Attacks: Yo-Yo Attack and Tandem Attack*. Red Hat. 17 March 2023.  
<https://www.youtube.com/watch?v=KTfY5NBIB8>
- ▶ [5] Kubernetes. Cluster Autoscaler. July 2023.  
<https://github.com/kubernetes/autoscaler/tree/master/cluster-autoscaler>